

Azure AI Services and Enterprise AI Systems

Executive Summary

Enterprise AI systems have shifted from isolated “model demos” to **production workloads** that must meet the same expectations as other business-critical cloud services: security, compliance, reliability, cost control, and operational visibility. Cloud platforms accelerate this shift by providing elastic infrastructure plus managed AI building blocks, so organizations can focus on solution design rather than running GPU clusters, building inference gateways, or maintaining search pipelines from scratch. Azure’s own guidance on AI workloads emphasizes that these systems often contain **non-deterministic behavior**, and therefore require special architectural attention across design, operations, and validation. ¹

Within Azure, the current AI platform experience increasingly converges around **Microsoft Foundry**, which is described by Microsoft ² as a unified Azure platform-as-a-service offering for enterprise AI operations, model builders, and application developers—aiming to reduce infrastructure management and centralize governance for AI apps and agents. ³ Complementing this, “Foundry Tools” provide ready-to-use AI APIs for modalities like Vision, Speech, Language, Document Intelligence, and Content Safety, with most capabilities available through REST APIs and SDKs. ⁴

For enterprise generative AI, Azure integrates large language models through **Azure OpenAI** as part of Foundry models, and supports retrieval grounding through **Azure AI Search**—a managed information retrieval platform positioned to connect enterprise data to agents and LLMs for grounded answers. ⁵ Critical enterprise controls are integrated into this ecosystem: identity via Microsoft Entra ID and managed identities (reducing reliance on long-lived secrets), centralized secret storage via Azure Key Vault, network isolation via private endpoints and disabling public network access, platform governance via Azure Policy, and compliance resources across global and industry frameworks. ⁶

Introduction

Cloud platforms are now the default substrate for modern AI systems because AI workloads demand a combination of resources and capabilities that are difficult to sustain on traditional fixed infrastructure: elastic compute (often including specialized accelerators), storage and data pipelines, secure connectivity to business systems, and continuous operations around monitoring, patching, and compliance. Azure’s Well-Architected Framework positions its guidance as a set of tenets and architectural decision points to help build a foundation for workloads on Azure, and AI-specific guidance highlights how AI workloads differ due to non-determinism and the tight coupling of code, data, and models. ⁷

A second reason cloud platforms matter is that they provide **managed platform services** that reduce “undifferentiated heavy lifting.” In AI, this includes not only GPU or distributed compute but also higher-level capabilities such as document OCR/understanding, semantic search and vector retrieval, speech transcription, translation, and content safety controls. Azure’s Architecture Center positions itself as a source of AI/ML guidance and links AI workload design to Well-Architected principles across pillars such as reliability, security, cost, and operational excellence. ⁸

Finally, enterprise AI systems increasingly need a governance and assurance posture that treats AI outputs as potentially risky operational artifacts. Even when an AI system is “just generating text,” it can influence customer interactions, employee decisions, and automated workflows. Azure’s AI workload guidance emphasizes validation and operational considerations for non-deterministic behavior—suggesting that testing, evaluation, and monitoring are architectural requirements rather than optional add-ons. ⁹

Overview of Azure AI services

Azure’s portfolio for building AI-powered applications can be understood as three complementary layers: (1) **prebuilt AI services** (“Foundry Tools”), (2) **model access and management** (“Foundry Models,” including Azure OpenAI), and (3) **end-to-end ML lifecycle** capabilities for custom models (Azure Machine Learning). This structure aligns with Microsoft’s messaging that Foundry is a unified platform for building, optimizing, and governing AI apps and agents at scale, while Foundry Tools provide modular capabilities via APIs and SDKs. ¹⁰

Foundry Tools: modality-first AI building blocks

Vision. Azure Vision provides algorithms for processing images and returning information based on visual features of interest. Microsoft’s documentation describes Azure Vision as cloud-based, allowing analysis of images via upload or URL and supporting a range of computer vision tasks. Image Analysis 4.0 is described as generally available with features such as synchronous OCR and people detection, while Microsoft also notes deprecation timelines for certain versions—an example of why enterprises must track service lifecycle as part of architecture governance. ¹¹

Speech. Azure Speech provides speech-to-text, text-to-speech, translation of spoken audio, and other speech capabilities. The service is positioned as a Foundry Tool accessed through a Foundry resource, supporting enterprise scenarios such as transcription and live voice interactions. ¹²

Language. Azure Language provides Natural Language Processing features for understanding and analyzing text, including capabilities such as entity extraction, intent classification, summarization, and question answering. Microsoft’s overview highlights that the service can be used through Foundry, REST APIs, and client libraries, and also notes integrations relevant to agent development (for example, tool endpoints and MCP-related integration). ¹³

Document Intelligence. Azure Document Intelligence is described as a cloud-based OCR and intelligent document processing service that extracts text, tables, structure, and key/value pairs from documents, and supports custom model training for structured and unstructured documents. This service often becomes the “front door” for enterprise document workflows (contracts, invoices, forms) that then feed analytics or downstream automation. ¹⁴

Content Safety. Azure AI Content Safety detects harmful user-generated and AI-generated content and provides text and image APIs to flag potentially harmful material. Microsoft also provides a Studio experience for exploration and testing, and separate documentation describing data privacy and security considerations for this service. ¹⁵

Retrieval as a first-class capability: Azure AI Search

Enterprise AI apps frequently depend on retrieval across internal knowledge bases, documents, and structured content. Azure AI Search is described as a fully managed, cloud-hosted service that “connects your data to AI,” unifying access to enterprise and web content so that agents and LLMs can use context and multi-source signals to produce grounded answers. Microsoft’s documentation explains that text is tokenized into inverted indexes and vectors are stored in vector indexes; AI enrichment can be used during indexing for chunking, vector generation, and structural transformations. ¹⁶

Azure AI Search also provides hybrid retrieval and semantic ranking, which are key to RAG-style grounding at enterprise quality. Hybrid search runs full-text and vector queries in parallel and merges results via Reciprocal Rank Fusion (RRF). Semantic ranker uses Microsoft language understanding models to rerank results and is described as integrated into “agentic retrieval.” ¹⁷

Model access and agent development: Foundry Models and model catalog

Microsoft positions Foundry Models as a destination for discovering, evaluating, and deploying AI models, and it differentiates between the “classic” portal and new portal experiences. Azure’s public model catalog emphasizes breadth of model providers and describes Azure OpenAI models as Microsoft-hosted OpenAI models with enterprise-grade security and compliance positioning. ¹⁸

Reference architecture for enterprise AI systems on Azure

A useful way to structure enterprise AI architecture is to separate the system into planes and layers: **experience layer** (channels and UI), **application layer** (business logic and orchestration), **AI capability layer** (prebuilt AI services and LLMs), **data/retrieval layer** (search, vector stores, documents), and **operations/governance plane** (identity, security, monitoring, compliance). Azure’s architecture guidance and AI workload guidance emphasize that AI workloads should follow Well-Architected principles, with special attention to non-determinism, data design, and operational aspects. ¹⁹

Core layers and typical Azure services

At a high level, a production enterprise AI system on Azure often uses:

- **Compute and hosting:** Azure App Service, Azure Kubernetes Service (AKS), Azure Functions, or other PaaS/container options for API hosting and orchestration. Azure Functions is described as a serverless solution to build robust apps with less infrastructure management, scaling automatically as needed. ²⁰
- **Data and content sources:** Azure Blob Storage, data platforms, SharePoint/OneLake (supported as sources for Azure AI Search) and enterprise systems-of-record. Azure AI Search documentation explicitly cites multiple supported data sources and describes indexed or remote access patterns based on freshness/latency/compliance needs. ¹⁶
- **AI services:** Vision, Speech, Language, Document Intelligence, Content Safety, and LLM endpoints through Azure OpenAI / Foundry Models, depending on the modality and task. ⁴
- **Retrieval and grounding:** Azure AI Search for indexing, vector retrieval, hybrid search, and semantic reranking to support RAG and agentic retrieval patterns. ²¹
- **Identity, secrets, and networking:** Microsoft Entra ID, managed identities, Azure Key Vault, private endpoints/Private Link, and disabling public network access for sensitive deployments. ²²

- **Observability:** Azure Monitor, Application Insights, OpenTelemetry-based tracing/logging/metrics to manage distributed AI apps and microservices, and to monitor quality-of-service. Azure Monitor describes the “three pillars of observability” (metrics, logs, distributed traces), and Application Insights is positioned as APM with OpenTelemetry integration. ²³

Architecture diagram: enterprise AI app with grounding and governance

```

flowchart LR
    U[Users / Channels] --> AP[App/API Layer<br/>App Service, AKS, Functions]
    AP --> ORCH[Orchestration & Business Logic<br/>workflows, agents, policy checks]
    ORCH --> LLM[LLM Endpoint<br/>Azure OpenAI / Foundry Models]
    ORCH --> AIS[AI Services APIs<br/>Vision, Speech, Language, Doc Intel, Safety]
    ORCH --> RET[Retrieval Layer<br/>Azure AI Search hybrid + semantic rerank]
    RET --> DS[Enterprise Data Sources<br/>Blob, DB, SharePoint, OneLake, etc.]

    subgraph GOV[Governance & Operations Plane]
        ID[Identity & Access<br/>Entra ID, Managed Identity]
        KV[Secrets & Keys<br/>Key Vault]
        NET[Network Controls<br/>Private endpoints, disable public access]
        MON[Telemetry & Monitoring<br/>Azure Monitor, Application Insights]
        POL[Policy & Compliance<br/>Azure Policy, compliance docs]
    end

    AP --- GOV
    ORCH --- GOV
    RET --- GOV
    LLM --- GOV
    AIS --- GOV

```

This diagram reflects a synthesis of Azure guidance: AI workloads need strong operational design; Azure AI Search is a managed retrieval platform that can ground agents and LLMs; and Azure supports enterprise controls (identity, secrets, networking, monitoring, policy) that must be treated as first-class architecture components rather than afterthoughts. ²⁴

Why “retrieval plus generation” is now a default enterprise pattern

Even when an organization starts with prebuilt AI APIs, enterprise AI apps usually converge on retrieval-driven designs because critical knowledge is proprietary, frequently updated, and permission-scoped. Azure AI Search documentation explains that classic search and more dynamic retrieval modes are complementary, and it emphasizes that indexing can include AI enrichment steps such as chunking and vector generation—capabilities that align directly with enterprise RAG architectures. ¹⁶

Azure OpenAI integration in enterprise applications

Azure integrates OpenAI-class models through “Azure OpenAI in Foundry Models,” positioned as providing access to cutting-edge models and emphasizing responsible use controls and limited access processes to ensure safeguards work in practice. ²⁵ In the Foundry documentation, Azure OpenAI

model availability varies by region, and Microsoft provides APIs (such as a Models List API) to enumerate available models for inference and fine-tuning on a given resource. ²⁶

Integration surface: endpoints, deployments, and “On Your Data”

A common enterprise integration approach is:

- 1) Create an Azure OpenAI resource and deploy relevant models (chat, reasoning, embeddings, multimodal, image generation depending on needs). ²⁷
- 2) Use the resource's inference endpoint from enterprise applications (web apps, APIs, internal tools).
- 3) For knowledge grounding, use Azure OpenAI “On Your Data,” which is described as enabling GPT models to run on enterprise data without training or fine-tuning, supporting more accurate responses by using designated data sources and exposing access via REST APIs, SDKs, or the Foundry portal UI. ²⁸
- 4) Implement retrieval via Azure AI Search (often configured as the retrieval backend), and apply best-practice controls such as managed identity authentication and private endpoints. ²⁹

From an architectural perspective, “On Your Data” is essentially a managed RAG pattern: the system retrieves supporting passages from specified sources and uses them to construct prompts for the model. Microsoft's documentation explicitly notes that you can guide the model to answer using retrieved documents only by specifying the behavior in the system message—illustrating that model behavior control is part of system design, not just model choice. ³⁰

Data privacy, abuse monitoring, and enterprise assurance

A core enterprise question is whether prompts and completions are used to train models. Microsoft's Azure OpenAI FAQ states that Azure OpenAI doesn't use customer data to retrain models and points to a data, privacy, and security guide; the Foundry “data privacy” documentation further describes how data provided to “Azure Direct Models” (including Azure OpenAI models) is processed, used, and stored to provide the service and for monitoring for uses that violate product terms. ³¹

This posture matters for enterprise governance because it defines the boundaries of data handling and informs decisions about what data can be sent to the model, how logs should be managed, and what additional controls are required under the organization's threat model and regulatory context. Azure also provides separate data privacy documentation for services like Content Safety, reinforcing that each AI component may have specific handling characteristics that must be part of architecture documentation. ³²

Network isolation: private endpoints and disabling public access

Azure documentation for networking configuration highlights patterns that are central to regulated deployments: disabling public network access, using private endpoints and private DNS, and controlling service-to-service access through managed identity and role-based access control when integrating Azure OpenAI with Azure AI Search and storage. ³³ The practical implication is that “enterprise” Azure OpenAI deployments often rely on a secure backbone architecture: application components within a virtual network (or connected via private endpoints) can access the AI service without exposing public endpoints, while identity and RBAC constrain who/what can call the service. ³⁴

AI development workflow on Azure

Enterprise AI development is no longer just “build a model, deploy a model.” It resembles a continuous lifecycle where teams must design for nondeterminism, evaluate behavior, monitor drift, manage costs, and respond to security and compliance requirements. Azure’s AI workload guidance frames these as architectural challenges, and Microsoft’s Well-Architected content highlights testing and evaluation practices as part of workload excellence. ⁹

Two parallel lifecycles: prebuilt AI + custom models

In Azure, many organizations operate two interlocking lifecycles:

- 1) **Prebuilt AI services lifecycle** (Vision, Speech, Language, Document Intelligence, Content Safety): integrate via REST APIs/SDKs, validate outputs against business criteria, monitor latency/costs, and manage service version changes and retirements. ³⁵
- 2) **Custom model lifecycle (MLOps)** for scenarios that require proprietary training, specialized predictions, or domain-specific performance. Azure Machine Learning documentation describes MLOps model management as applying DevOps principles (CI/CD) to ML, supporting reproducible pipelines, reusable environments, model registration/versioning, lineage tracking, monitoring, alerts, and end-to-end lifecycle automation. ³⁶

This duality is a strength of Azure’s platform approach: teams can begin with prebuilt services for rapid time-to-value and extend with custom ML when needed, while using platform services to manage operational complexity. ³⁷

GenAI development and evaluation: prompt flow and evaluation runs

For generative AI apps, development often centers on prompts, retrieval strategies, tool orchestration, and systematic evaluation rather than training from scratch. Prompt flow is described as a development tool that streamlines the development cycle of AI applications powered by LLMs, supporting prototyping, iteration, and deployment. ³⁸

Evaluation is increasingly treated as a first-class production concern. Foundry documentation describes how to create evaluation runs against models, agents, or test datasets using built-in metrics, and it supports custom evaluation flows for flexibility. ³⁹ This matters because nondeterministic outputs require continuous validation: regression testing for prompts, retrieval changes, safety policies, and model updates becomes analogous to how enterprises perform reliability engineering for microservices.

Observability for enterprise AI apps: traces, logs, and metrics

In production, enterprise AI systems behave like distributed systems: they comprise an application layer plus calls to multiple AI services, search indexes, and downstream systems-of-record. Azure Monitor describes the three pillars of observability—metrics, logs, and distributed traces—and emphasizes correlating data across pillars for sufficient observability. ⁴⁰ Application Insights provides APM and supports OpenTelemetry-based telemetry collection, enabling organizations to trace request paths and performance across services. ⁴¹

A practical best practice is to treat each AI capability call (LLM inference, retrieval query, document extraction) as a traceable operation with correlation IDs, so that engineers can diagnose failures,

latency spikes, cost anomalies, or unexpected output changes. Microsoft’s documentation explicitly highlights approaches to collect and manage telemetry with OpenTelemetry and to integrate Application Insights into broader monitoring, reinforcing that observability is part of the expected enterprise setup. ⁴²

Enterprise use cases and solution patterns

Azure’s AI services portfolio maps cleanly to common enterprise use cases because these use cases are fundamentally modality-driven: language understanding for customer interactions, document intake for operations, and retrieval for knowledge work. Sample applications, tutorials, and architecture guidance provide concrete patterns that can be generalized.

Intelligent chatbots and copilots grounded on enterprise content

A common enterprise pattern is a chatbot that answers questions about internal policies, product docs, or operational runbooks. The widely referenced Azure sample “chat with Azure OpenAI and Azure AI Search” demonstrates a ChatGPT-like experience over “your own data,” illustrating RAG grounding in an enterprise-friendly architecture. ⁴³

In more managed patterns, Azure OpenAI “On Your Data” is presented as a way to run GPT models on enterprise data without training or fine-tuning, specifying sources for more accurate responses based on the latest information available in designated data sources. ²⁸ Retrieval quality can be increased by using Azure AI Search hybrid search and semantic ranker; hybrid retrieval merges vector and full-text results via RRF and semantic ranker reranks results to improve relevance, which is directly aligned with enterprise goals of “high precision and strong grounding.” ¹⁷

Document processing and workflow automation

Many organizations have “document bottlenecks” in finance, legal, healthcare, procurement, and HR. Azure Document Intelligence is designed to extract text, tables, structure, and key/value pairs from documents and to support both prebuilt and custom models, enabling downstream validation and automation. ¹⁴

In enterprise architectures, document extraction rarely stands alone. It often feeds: - a search index (for discovery and retrieval), - business rules engines (for validations), - analytics platforms (for dashboards and audits), - and workflow systems (for approvals and case management).

Azure AI Search documentation notes AI enrichment during indexing—such as chunking text and transforming content into searchable JSON—suggesting a natural integration between document extraction and retrieval layers. ¹⁶

Automation and event-driven “AI in the loop”

Automation use cases often benefit from serverless and event-driven patterns. Azure Functions is described as a serverless solution allowing code to run without managing servers, and it is commonly used to glue together event triggers, message queues, and API calls. ²⁰ In an AI system, Functions can be used to trigger document extraction when files arrive, call Content Safety before publishing content, or initiate retrieval updates when new documentation is deployed.

The enterprise benefit is architectural: by embedding AI capabilities inside event-driven workflows, organizations can automate repetitive processes while introducing AI only at the “decision points” that require language understanding, vision analysis, or generation. This reduces cost and risk compared with making the entire workflow agentic.

Analytics and decision support

In decision support settings, AI services often appear as enrichment layers for analytics: extracting structured features from text, classifying documents, detecting entities, summarizing trends, and then feeding BI or data warehouses. Azure Machine Learning’s MLOps documentation emphasizes reproducible pipelines, model tracking, and monitoring, which are relevant for predictive analytics and operational decision models. ⁴⁴

Security, compliance, and responsible AI

Enterprise AI systems require security and governance controls that account for both traditional cloud risks and AI-specific risks (such as unintended disclosure through prompts, unsafe content generation, or emergent behavior in non-deterministic systems). Microsoft provides multiple layers of guidance: responsible use documentation for Foundry Tools, Responsible AI Standard artifacts, compliance documentation, and platform governance/security services.

Identity and access: Entra ID, managed identities, least privilege

A widely adopted enterprise pattern on Azure is to minimize long-lived secrets by relying on identity-based authentication. Microsoft’s managed identities documentation explains that managed identities allow authentication to resources that support Microsoft Entra authentication without managing credentials, and that credentials are not accessible to developers. ⁴⁵ This aligns with a least-privilege architecture for AI systems where each component (app, ingestion job, retrieval indexer) has its own identity and only the minimum roles needed to operate.

Secrets, keys, and cryptography: Azure Key Vault

AI systems often require secure handling of API keys, TLS certificates, encryption keys, and connection secrets. Azure Key Vault is described as a secure secrets store that enables apps to store and use keys, secrets, and certificates (“objects”), reducing accidental leakage and supporting RBAC/access policies, encryption at rest backed by hardware security modules, and monitoring of access via logging. ⁴⁶

In AI architectures, Key Vault supports secure integration to AI services, storage accounts, and downstream systems—especially where regulatory controls require strong key custody and auditing.

Network security: private endpoints and controlled exposure

For sensitive workloads, network isolation is frequently mandatory. Azure documentation for configuring Azure OpenAI networking describes creating private endpoints and private DNS zones and disabling public network access so that private endpoint connections become the exclusive access path. ⁴⁷ Similar patterns exist for Azure AI Search private endpoints, and documentation explicitly describes scenarios where Azure OpenAI and Azure AI Search are connected via private endpoints for “use your own data” configurations. ⁴⁸

This is significant for enterprise AI risk management: a private endpoint posture reduces exposure of AI and retrieval services to the public internet and supports internal-only traffic patterns over Azure backbone networks.

Governance and compliance controls: Azure Policy, compliance offerings, and Purview

Azure Policy is described as enforcing organizational standards and assessing compliance at scale, with a compliance dashboard and remediation capabilities for existing and new resources. ⁴⁹ From an AI system perspective, Azure Policy can enforce baseline configuration standards: resource locations (data residency), encryption settings, public network access disabled, tagging requirements for ownership and data classification, and other guardrails.

For regulatory alignment, Azure's compliance documentation provides a central entry point for compliance in Azure and links to compliance offerings across different regimes and standards. ⁵⁰ This portfolio matters for enterprise AI systems because AI workloads frequently handle sensitive data and must meet industry or regional obligations; compliance documentation becomes part of the evidence chain for audits and risk assessments.

For organizations using Microsoft 365 copilots, agents, or broader AI usage monitoring, Microsoft Purview provides capabilities positioned as helping mitigate and manage risks associated with AI usage, including Data Security Posture Management for AI (DSPM for AI). Microsoft's Purview documentation describes using Purview to implement additional governance controls and monitor AI use across the enterprise. ⁵¹ While Purview is not limited to Azure AI Services, it is relevant to enterprise AI governance because many organizations operate AI across both Azure-hosted apps and broader productivity ecosystems.

Responsible AI: principles, standards, and service-level safeguards

Microsoft provides a "Responsible use of AI" overview for Foundry Tools, pointing to service-specific guidance for generative AI and modality services like Vision and Video Indexer. ⁵² At the organizational level, Microsoft publishes its Responsible AI Standard v2 general requirements as a formal artifact, and it provides an Impact Assessment Guide intended to support teams conducting responsible AI impact assessments required by the standard. ⁵³

At the service level, Azure AI Content Safety provides explicit safeguards against harmful content via detection APIs and Studio tooling, and its documentation emphasizes that the customer remains responsible for implementation decisions even when the service is designed with privacy and security in mind. ¹⁵ The practical enterprise implication is that responsible AI is not only a policy document—it becomes a set of engineering controls embedded in the architecture: safety filtering, logging and escalation policies, human review for sensitive scenarios, and evaluation against responsible AI requirements.

Benefits and closing synthesis

Azure enables enterprise AI systems by combining (1) modular AI capabilities (Vision, Speech, Language, Document Intelligence, Content Safety), (2) a managed retrieval platform for grounding (Azure AI Search), (3) model access and app/agent tooling (Foundry), and (4) enterprise platform controls (identity, secrets, networking, observability, policy, compliance documentation). The "cloud advantage" here is not simply scalability; it is the **integration of AI capabilities into a governed, operable platform** where security and compliance patterns are supported by design, and where

lifecycle concerns (monitoring, evaluation, and service retirement awareness) can be handled systematically. ⁵⁴

References

Core platform and AI workload architecture

- Azure Well-Architected Framework (overview): <https://learn.microsoft.com/en-us/azure/well-architected/>
- AI workloads on Azure (Well-Architected): <https://learn.microsoft.com/en-us/azure/well-architected/ai/get-started>
- Azure Architecture Center (AI/ML): <https://learn.microsoft.com/en-us/azure/architecture/ai-ml/>
- Cloud Adoption Framework – AI PaaS architectures: <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/scenarios/ai/platform/architectures>

Microsoft Foundry, Foundry Tools, and model access

- Microsoft Foundry documentation hub: <https://learn.microsoft.com/en-us/azure/foundry/>
- What is Microsoft Foundry?: <https://learn.microsoft.com/en-us/azure/foundry/what-is-foundry>
- What are AI services / Foundry Tools?: <https://learn.microsoft.com/en-us/azure/ai-services/what-are-ai-services>
- Foundry portal: <https://ai.azure.com/>
- Model catalog: <https://ai.azure.com/catalog>

Azure AI Services (Foundry Tools)

- Azure Vision documentation: <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/>
- Image Analysis overview: <https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/overview-image-analysis>
- Azure Speech overview: <https://learn.microsoft.com/en-us/azure/ai-services/speech-service/overview>
- Azure Language documentation: <https://learn.microsoft.com/en-us/azure/ai-services/language-service/>
- Azure Language overview: <https://learn.microsoft.com/en-us/azure/ai-services/language-service/overview>
- Document Intelligence overview: <https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/overview?view=doc-intel-4.0.0>
- Azure AI Content Safety overview: <https://learn.microsoft.com/en-us/azure/ai-services/content-safety/overview>

Azure AI Search (retrieval, RAG, and semantic ranking)

- Azure AI Search overview: <https://learn.microsoft.com/en-us/azure/search/search-what-is-azure-search>
- Hybrid search overview: <https://learn.microsoft.com/en-us/azure/search/hybrid-search-overview>
- Hybrid search ranking (RRF): <https://learn.microsoft.com/en-us/azure/search/hybrid-search-ranking>

- Semantic search (semantic ranker) overview: <https://learn.microsoft.com/en-us/azure/search/semantic-search-overview>
- Azure AI Search features list: <https://learn.microsoft.com/en-us/azure/search/search-features-list>

Azure OpenAI / Foundry OpenAI integration and privacy/security

- Azure OpenAI models product page: <https://azure.microsoft.com/en-us/products/ai-foundry/models/openai>
- Working with Azure OpenAI models in Foundry: <https://learn.microsoft.com/en-us/azure/foundry/openai/how-to/working-with-models>
- Azure OpenAI FAQ (classic): <https://learn.microsoft.com/en-us/azure/foundry-classic/openai/faq>
- Azure OpenAI On Your Data (classic): <https://learn.microsoft.com/en-us/azure/foundry-classic/openai/concepts/use-your-data>
- Network and access configuration for On Your Data: <https://learn.microsoft.com/en-us/azure/foundry-classic/openai/how-to/on-your-data-configuration>
- Configure Azure OpenAI networking: <https://learn.microsoft.com/en-us/azure/foundry-classic/openai/how-to/network>
- Data, privacy, and security for Azure Direct Models (includes Azure OpenAI models): <https://learn.microsoft.com/en-us/azure/foundry/responsible-ai/openai/data-privacy>

MLOps and custom model lifecycle

- MLOps model management and deployment (Azure ML): <https://learn.microsoft.com/en-us/azure/machine-learning/concept-model-management-and-deployment?view=azureml-api-2>
- Azure ML serverless compute for training: <https://learn.microsoft.com/en-us/azure/machine-learning/how-to-use-serverless-compute?view=azureml-api-2>

Security, identity, monitoring, and governance

- Managed identities overview: <https://learn.microsoft.com/en-us/entra/identity/managed-identities-azure-resources/overview>
- Microsoft Entra ID documentation: <https://learn.microsoft.com/en-us/entra/identity/>
- Azure Key Vault overview: <https://learn.microsoft.com/en-us/azure/key-vault/general/overview>
- Azure Policy overview: <https://learn.microsoft.com/en-us/azure/governance/policy/overview>
- Azure compliance documentation: <https://learn.microsoft.com/en-us/azure/compliance/>
- Compliance offerings overview: <https://learn.microsoft.com/en-us/azure/compliance/offerings/>

Responsible AI

- Responsible use of AI (Foundry Tools): <https://learn.microsoft.com/en-us/azure/ai-services/responsible-use-of-ai-overview>
- Microsoft Responsible AI Standard v2 (general requirements PDF): <https://cdn-dynmedia-1.microsoft.com/is/content/microsoftcorp/microsoft/final/en-us/microsoft-brand/documents/Microsoft-Responsible-AI-Standard-General-Requirements.pdf>

- Responsible AI Impact Assessment Guide (PDF): <https://msblogs.thesourcemediaassets.com/sites/5/2022/06/Microsoft-RAI-Impact-Assessment-Guide.pdf>

-
- 1 9 24 **AI workloads on Azure - Well-Architected Framework**
https://learn.microsoft.com/en-us/azure/well-architected/ai/get-started?utm_source=chatgpt.com
 - 2 46 **Azure Key Vault Overview**
https://learn.microsoft.com/en-us/azure/key-vault/general/overview?utm_source=chatgpt.com
 - 3 10 54 **What is Microsoft Foundry? - Microsoft ...**
https://learn.microsoft.com/en-us/azure/foundry/what-is-foundry?utm_source=chatgpt.com
 - 4 35 **What are Foundry Tools? - Foundry Tools**
https://learn.microsoft.com/en-us/azure/ai-services/what-are-ai-services?utm_source=chatgpt.com
 - 5 25 27 **Azure OpenAI in Foundry Models**
https://azure.microsoft.com/en-us/products/ai-foundry/models/openai?utm_source=chatgpt.com
 - 6 22 45 **Managed identities for Azure resources**
https://learn.microsoft.com/en-us/entra/identity/managed-identities-azure-resources/overview?utm_source=chatgpt.com
 - 7 **Azure Well-Architected Framework**
https://learn.microsoft.com/en-us/azure/well-architected/?utm_source=chatgpt.com
 - 8 19 **AI Architecture Design - Azure Architecture Center**
https://learn.microsoft.com/en-us/azure/architecture/ai-ml/?utm_source=chatgpt.com
 - 11 **Azure Vision in Foundry Tools documentation**
https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/?utm_source=chatgpt.com
 - 12 **What Is Azure Speech? - Foundry Tools**
https://learn.microsoft.com/en-us/azure/ai-services/speech-service/overview?utm_source=chatgpt.com
 - 13 **What is Azure Language in Foundry Tools**
https://learn.microsoft.com/en-us/azure/ai-services/language-service/overview?utm_source=chatgpt.com
 - 14 **What Is Azure Document Intelligence in Foundry Tools?**
https://learn.microsoft.com/en-us/azure/ai-services/document-intelligence/overview?view=doc-intel-4.0.0&utm_source=chatgpt.com
 - 15 **What is Azure AI Content Safety?**
https://learn.microsoft.com/en-us/azure/ai-services/content-safety/overview?utm_source=chatgpt.com
 - 16 21 43 **Introduction to Azure AI Search**
https://learn.microsoft.com/en-us/azure/search/search-what-is-azure-search?utm_source=chatgpt.com
 - 17 **Hybrid Search - Azure AI Search**
https://learn.microsoft.com/en-us/azure/search/hybrid-search-overview?utm_source=chatgpt.com
 - 18 **Microsoft Foundry Models overview (classic)**
https://learn.microsoft.com/en-us/azure/foundry-classic/concepts/foundry-models-overview?utm_source=chatgpt.com
 - 20 **Azure Functions overview**
https://learn.microsoft.com/en-us/azure/azure-functions/functions-overview?utm_source=chatgpt.com
 - 23 40 **Azure Monitor data platform**
https://learn.microsoft.com/en-us/azure/azure-monitor/fundamentals/data-platform?utm_source=chatgpt.com

- 26 **Azure OpenAI in Microsoft Foundry Models working with ...**
https://learn.microsoft.com/en-us/azure/foundry/openai/how-to/working-with-models?utm_source=chatgpt.com
- 28 30 **Azure OpenAI On Your Data (classic)**
https://learn.microsoft.com/en-us/azure/foundry-classic/openai/concepts/use-your-data?utm_source=chatgpt.com
- 29 **Network and access configuration for Azure OpenAI On ...**
https://learn.microsoft.com/en-us/azure/foundry-classic/openai/how-to/on-your-data-configuration?utm_source=chatgpt.com
- 31 **Azure OpenAI frequently asked questions**
https://learn.microsoft.com/en-us/azure/foundry-classic/openai/faq?utm_source=chatgpt.com
- 32 **Data, privacy, and security for Azure AI Content Safety**
https://learn.microsoft.com/en-us/azure/foundry/responsible-ai/content-safety/data-privacy?utm_source=chatgpt.com
- 33 34 47 **Configure Azure OpenAI networking (classic)**
https://learn.microsoft.com/en-us/azure/foundry-classic/openai/how-to/network?utm_source=chatgpt.com
- 36 37 44 **MLOps model management with Azure Machine Learning**
https://learn.microsoft.com/en-us/azure/machine-learning/concept-model-management-and-deployment?view=azureml-api-2&utm_source=chatgpt.com
- 38 **Prompt flow in Microsoft Foundry portal (classic)**
https://learn.microsoft.com/en-us/azure/foundry-classic/concepts/prompt-flow?utm_source=chatgpt.com
- 39 **Evaluate generative AI models and applications ...**
https://learn.microsoft.com/en-us/azure/foundry/how-to/evaluate-generative-ai-app?utm_source=chatgpt.com
- 41 **Introdução ao Application Insights - Azure Monitor**
https://learn.microsoft.com/pt-br/azure/azure-monitor/app/app-insights-overview?utm_source=chatgpt.com
- 42 **Collect OpenTelemetry (OTel) for Application Insights ...**
https://learn.microsoft.com/en-us/azure/azure-monitor/app/opentelemetry-overview?utm_source=chatgpt.com
- 48 **Create a private endpoint for a secure connection**
https://docs.azure.cn/en-us/search/service-create-private-endpoint?utm_source=chatgpt.com
- 49 **Overview of Azure Policy**
https://learn.microsoft.com/en-us/azure/governance/policy/overview?utm_source=chatgpt.com
- 50 **Azure compliance documentation**
https://learn.microsoft.com/en-us/azure/compliance/?utm_source=chatgpt.com
- 51 **Microsoft Purview data security and compliance protections ...**
https://learn.microsoft.com/en-us/purview/ai-microsoft-purview?utm_source=chatgpt.com
- 52 **Overview of Responsible use of AI - Foundry Tools**
https://learn.microsoft.com/en-us/azure/ai-services/responsible-use-of-ai-overview?utm_source=chatgpt.com
- 53 **Microsoft Responsible AI Standard v2 General Requirements**
https://cdn-dynmedia-1.microsoft.com/is/content/microsoftcorp/microsoft/final/en-us/microsoft-brand/documents/Microsoft-Responsible-AI-Standard-General-Requirements.pdf?utm_source=chatgpt.com